

Deep Learning Practical Guidance

Getting Started with Image & Text

Amit Kapoor [@amitkaps](#)

Bargava Subramanian [@bargava](#)


Anand Chitipothu [@anandology](#)

Bootcamp Approach


- **Domain:** Image & Text
- **Applied:** Proven & Practical
- **Intuition:** Visualisation & Analogies
- **Code:** Learning by Doing
- **Math:** Attend HackerMath!

Learning Paradigm


Classical Programming Paradigm



Learning Paradigm - Machine Learning



Learning Paradigm - Deep Learning



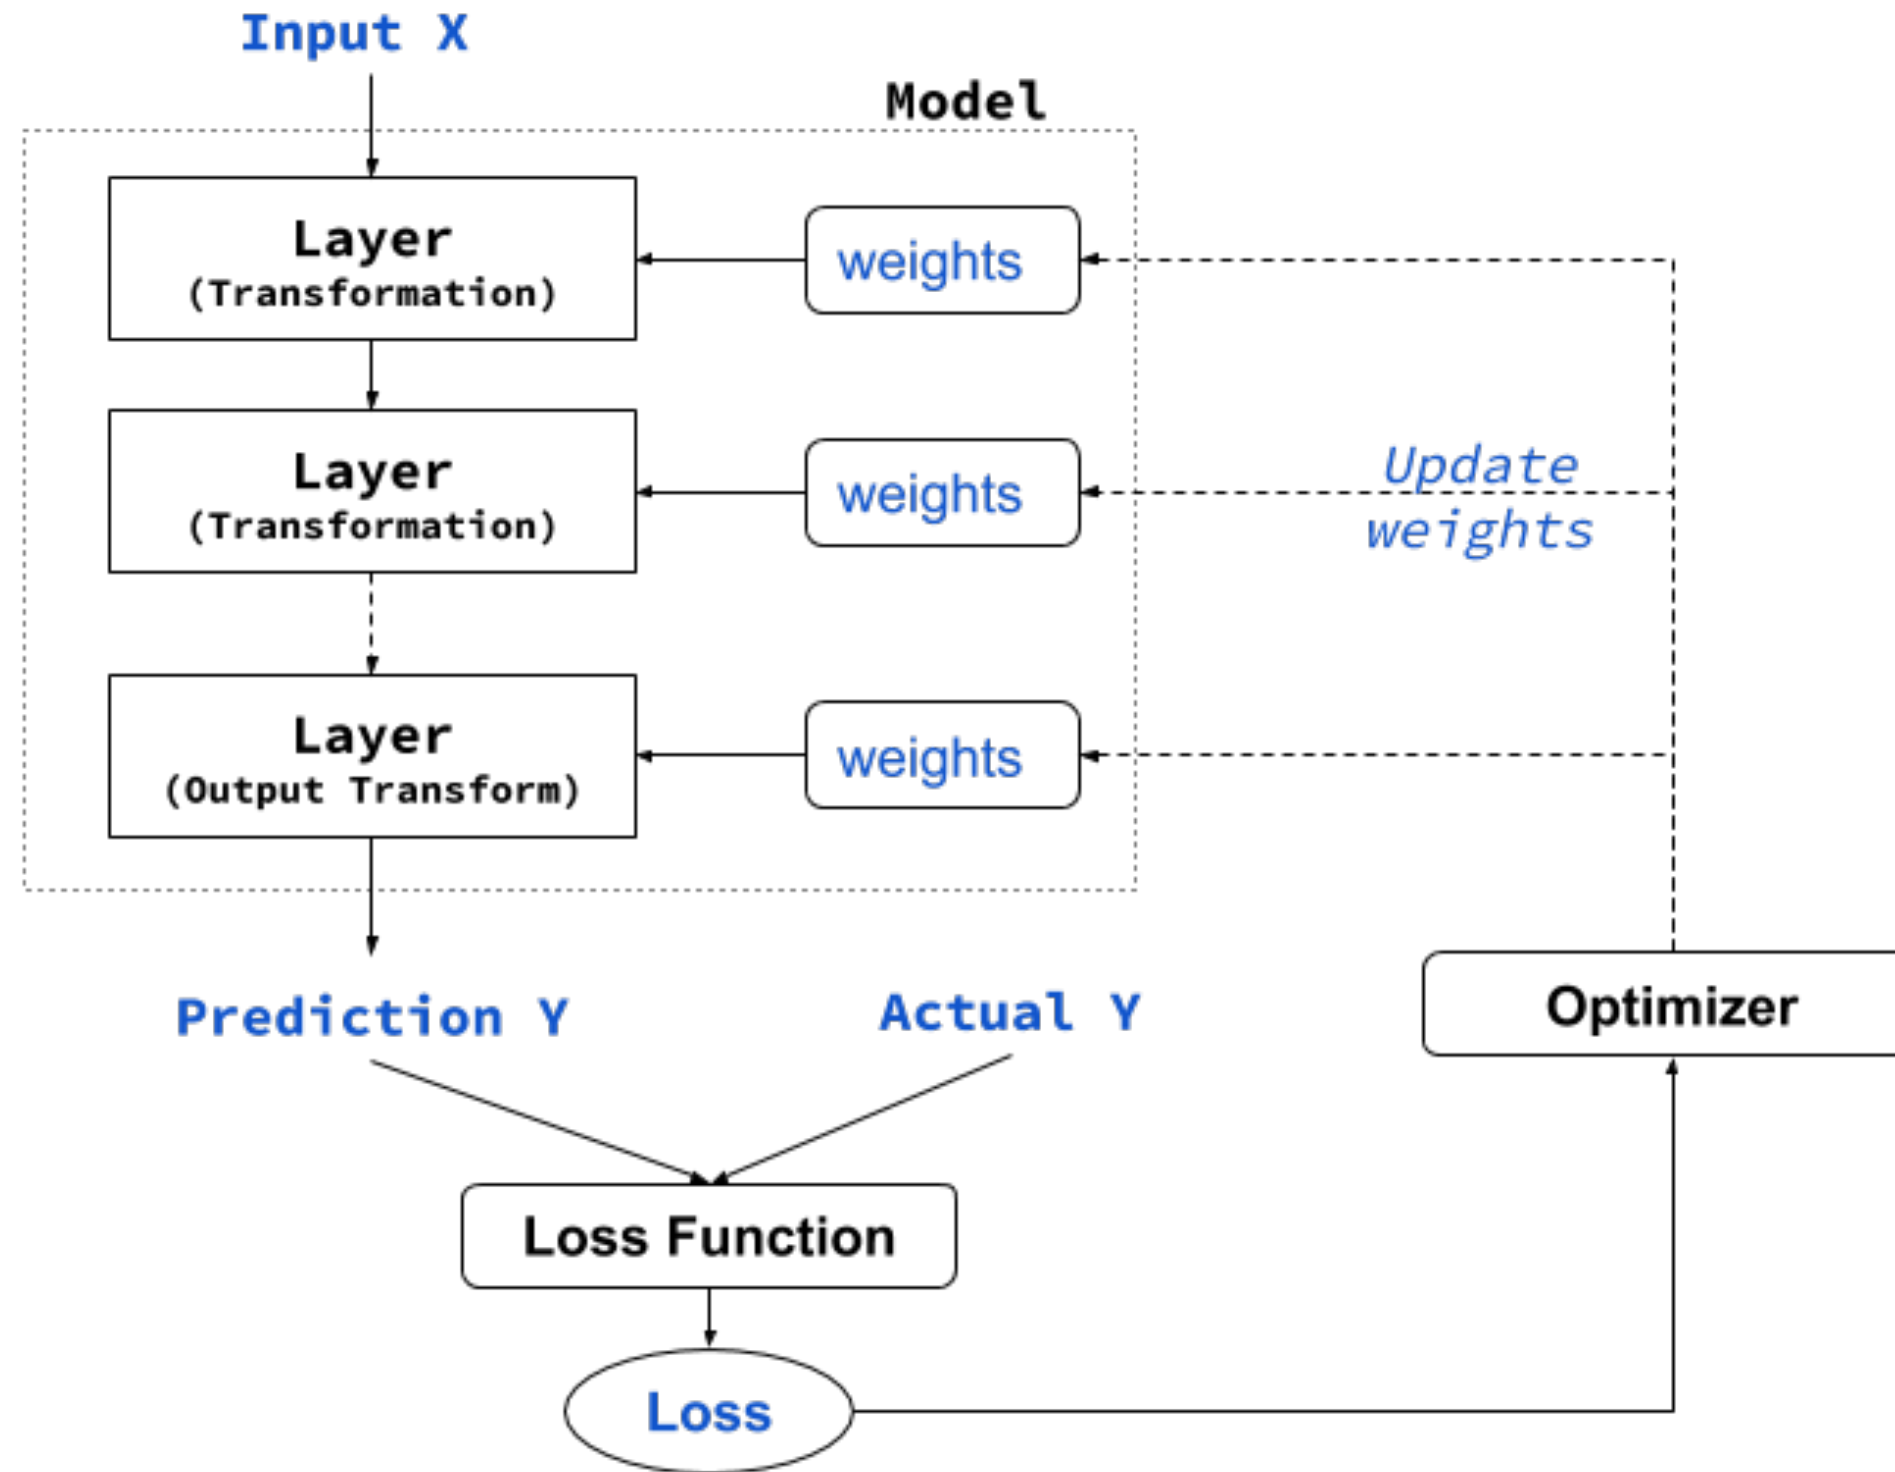
Learning Types & Applications

- **Supervised:** Regression, Classification, ...
- **Unsupervised:** Dimensionality Reduction, Clustering, ...
- **Self (Semi)-supervised:** Auto-encoders, Generative Adversarial Network, ...
- **Reinforcement Learning:** Games, Self-Driving Car, Robotics, ...

Focus: Supervised Learning

- **Classification:** Image, Text, Speech, Translation
- Sequence generation: Given a picture, predict a caption describing it.
- Syntax tree prediction: Given a sentence, predict its decomposition into a syntax tree.
- Object detection: Given a picture, draw a bounding box around certain objects inside the picture.
- Image segmentation: Given a picture, draw a pixel-level mask on a specific object.

Learning Approach



Data Representation: Tensors

- Numpy arrays (aka Tensors)
- Generalised form of matrix (2D array)
- Attributes
 - Axes or Rank: `ndim`
 - Dimensions: `shape` e.g. `(5, 3)`
 - Data Type: `dtype` e.g. `float32`, `uint8`, `float64`

Tensor Types

- **Scalar:** 0D Tensor
- **Vector:** 1D Tensor
- **Matrix:** 2D Tensor
- **Higher-order:** 3D, 4D or 5D Tensor

Input X


Tensor	Example	Shape
2D	Tabular	(samples, features)
3D	Sequence	(samples, steps, features)
4D	Images	(samples, height, width, channels)
5D	Videos	(samples, frames, height, width, channels)

Learning Unit

$$y = \text{RELU}(\text{dot}(w, x) + \text{bias})$$

weights are $w_1 \dots w_n$ & activation is RELU

$$f(z) = \max(z, 0)$$



Model Architecture

Basic Model: **Sequential** - A linear stack of layers.

Core Layers

- Dense Layers: Fully connected layer of learning units
(Also called Multi-layer Perceptron)
- Flatten

Output y & Loss

y	Last Layer Activation	Loss Function
Binary Class	sigmoid	Binary Crossentropy
Multi Class	softmax	Categorical Crossentropy
Multi Class Multi Label	sigmoid	Binary Crossentropy
Regression	None	Mean Square Error
Regression (0-1)	sigmoid	MSE or Binary Crossentropy

Optimizers

- **SGD**: Excellent but requires tuning learning-rate decay, and momentum parameters
- **RMSProp**: Good for RNNs
- **Adam**: Adaptive momentum optimiser, generally a good starting point.

Guidance for DL

*General guidance on building and training neural networks.
Treat them as heuristics (derived from experimentation) and
as good starting points for your own explorations.*

Pre-Processing

- **Normalize / Whiten** your data (Not for text!)
- **Scale** your data appropriately (for outlier)
- **Handle Missing Values** - Make them 0 (Ensure it exists in training)
- **Create Training & Validation Split**
- **Stratified** split for multi-class data
- **Shuffle** data for non-sequence data. Careful for sequence!!

General Architecture

- Use **ADAM** Optimizer (to start with)
- Use **RELU** for non-linear activation (Faster for learning than others)
- Add **Bias** to each layer
- Use **Xavier** or **Variance-Scaling** initialisation (Better than random initialisation)
- Refer to output layers activation & loss function guidance for tasks

Dense / MLP Architecture

- No. of units reduce in deeper layer
- Units are typically 2^n
- Don't use more than 4 - 5 layers in dense networks

CNN Architecture (for Images)

- Increase **Convoluton** filters as you go deeper from 32 to 64 or 128 (Max)
- Use **Pooling** to subsample: Makes image robust from translation, scaling, rotation
- Use **pre-trained models** as *feature extractors* for similar tasks
- Progressively **train n-last layers** if the model is not learning
- **Image Augmentation** is key for small data and for

RNN / CNN Architecture (for NLP)

- **Embedding layer is critical. Words are better than Characters**
- Learn the embedding with the task or use pre-trained embedding as starting point
- Use BiLSTM / LSTM vs Simple RNN. Remember, RNNs are really slow to train
- Experiment with 1D CNN with larger kernel size (7 or 9) than used for images.
- MLP can work with bi-grams for many simple tasks.

Learning Process

- **Validation Process**
 - Large Data: Hold-Out Validation
 - Smaller Data: K-Fold (Stratified) Validation
- **For Underfitting**
 - Add more layers: **go Deeper**
 - Make the layers bigger: **go wider**
 - Train for more epochs

Learning Process

- **For Overfitting**
 - Get more training data (e.g. actual or image augmentation)
 - Reduce Model Capacity
 - Add weight regularisation (e.g. L1, L2)
 - Add Dropouts or use Batch Normalization